

# Chapter 4 Hypothesis testing

## Example 4.3 Two samples

Test the difference of two means. The box-plot is not helpful.


```
[34]  rng = 123;
      x = normrnd(0, 1, 26,1);
      y = normrnd(0.5, 0.8, 15,1);
      x'
      y'

      grp = [ones(1, 26), ones(1, 15)*2];
      grp
      boxplot([x;y], grp)
```

```
ans =
  Columns 1 through 7
    0.5377    1.8339   -2.2588    0.8622    0.3188   -1.3077   -0.4336
  Columns 8 through 14
    0.3426    3.5784    2.7694   -1.3499    3.0349    0.7254   -0.0631
  Columns 15 through 21
    0.7147   -0.2050   -0.1241    1.4897    1.4090    1.4172    0.6715
  Columns 22 through 26
   -1.2075    0.7172    1.6302    0.4889    1.0347

ans =
  Columns 1 through 7
    1.0815    0.2572    0.7351   -0.1298    1.2107   -0.4177   -0.3551
  Columns 8 through 14
   -0.1476   -1.8554    1.6507    0.7602   -0.1039    1.5962   -0.8692
  Column 15
    0.4182

grp =
  Columns 1 through 13
    1    1    1    1    1    1    1    1    1    1    1    1
1
  Columns 14 through 26
    1    1    1    1    1    1    1    1    1    1    1    1
1
  Columns 27 through 39
    2    2    2    2    2    2    2    2    2    2    2    2
2
  Columns 40 through 41
    2    2
```



### Warning on tail probabilities

Calculation of upper tail probabilities of a discrete distribution in Matlab.

Let  $X \sim \text{Bin}(25, 0.6)$ .

```
[35] n = 25; p = 0.6;
```

Example 1. Calculate  $P(X \leq 15)$ .

```
[36] sum(binopdf(0:15, n,p))
```

```
ans =  
    0.5754
```

```
[37] binocdf(15, n, p)
```

```
ans =  
    0.5754
```

That's easy.

Example 2. Now Calculate  $P(X \geq 20)$ .

```
[38] 1- binocdf(20, n,p)
      disp('WRONG!')
```

```
ans =
      0.0095
WRONG!
```

```
[39] binocdf(20, n,p, 'upper')
      disp('WRONG!')
```

```
ans =
      0.0095
WRONG!
```

```
[40] sum(binopdf(20:25, n,p))
      disp('CORRECT.')
```

```
ans =
      0.0294
CORRECT.
```

So you must be careful!

You can do it correctly with `cdfbino` using  $x = 20 - 1$ :

```
[41] binocdf(19, n,p, 'upper')
      1- binocdf(19, n,p)
```


```
ans =
      0.0294
ans =
      0.0294
```

Power function  $p \mapsto P_p(X \geq 59)$ .

```
[42] pow_bin = @(p) binocdf(58, 100, p, 'upper');
```

```
[43] fplot(pow_bin, [0,1])  
vline(0.5, '-')
```

```
hline(0.05)
```




**Figure 4.4**

The function  $x \mapsto P_{0.5}(X \geq x)$  for  $X \sim \text{Bin}(100, 0.5)$ .


```
[44] bino = @(x) 1 - binocdf(x-1, 100, 0.5);
```

```
[45] fplot(bino, [0, 100])  
  
hline(0.05)
```



```
[46] fplot(bino, [57, 60])
```


```
hline(0.044)
hline(0.05)
hline(0.067)
```



**Figure 4.5 Power function for the Gauss test**

- Red. One sided test  $H_0 : \mu \leq 0$  vs  $H_1 : \mu > 0$
- Blue. One sided test  $H_0 : \mu \geq 0$  vs  $H_1 : \mu < 0$
- Yellow. Two-sided test  $H_0 : \mu = \mu_0$  vs  $H_1 : \mu \neq \mu_0$

```
[47] n = 5;  
a0 = 0.05;  
mu0 = 0;  
sigma = 2;  
pow_gauss_R = @(mu) 1 - normcdf(norminv(1-a0,0,1) + sqrt(n) .* (mu0 - mu),0,1);  
fplot(pow_gauss_R, [-3,3])  
  
hold on  
  
pow_gauss_L = @(mu) normcdf(-norminv(1-a0,0,1) + sqrt(n) .* (mu0 - mu),0,1);  
fplot(pow_gauss_L, [-3,3])  
  
pow_gauss_2 = @(mu) 1 - normcdf( norminv(1-a0/2) - sqrt(n) .* (mu - mu0),0,1) +  
normcdf( norminv(1-a0/2) + sqrt(n) .* (mu - mu0),0,1);  
fplot(pow_gauss_2, [-3,3])  
  
hline(0.05)  
hold off
```




Note that `normcdf(x)` is the distribution function of  $N(0, 1)$  that is  $\Phi(x)$ , and `norminv(a)` is the  $\alpha$ -quantile of the  $N(0, 1)$ .

### Example 4.15 Gauss test, power function

Compare  $n = 5$  with  $n = 50$ .

```
[48] a0 = 0.05;
      mu0 = 0;
      sigma = 2;
      n = 5;
      pow_gauss_R = @(mu) 1 - normcdf(norminv(1-a0,0,1) + sqrt(n) .* (mu0 - m
      fplot(pow_gauss_R, [-3,3])
      hold on
      n = 50;
      pow_gauss_R = @(mu) 1 - normcdf(norminv(1-a0,0,1) + sqrt(n) .* (mu0 - m
      fplot(pow_gauss_R, [-3,3])
      hline(a0)
      vline(0, '-k')
      hold off
```




Compare  $\sigma = 2$  with  $\sigma = 4$ .

```

[49] n = 50;
a0 = 0.05;
mu0 = 0;
sigma = 2;
pow_gauss_R = @(mu) 1 - normcdf(norminv(1-a0,0,1) + sqrt(n) .* (mu0 - m
fplot(pow_gauss_R, [-3,3])
hold on
sigma = 4;
pow_gauss_R = @(mu) 1 - normcdf(norminv(1-a0,0,1) + sqrt(n) .* (mu0 - m
fplot(pow_gauss_R, [-3,3])
hline(a0)
vline(0,'-k')
hold off

```



Compare  $\alpha_0 = 0.05$  with  $\alpha_0 = 0.10$ .

```

[50] n = 10;
a0 = 0.05;
mu0 = 0;
sigma = 2;
pow_gauss_R = @(mu) 1 - normcdf(norminv(1-a0,0,1) + sqrt(n) .* (mu0 - m
fplot(pow_gauss_R, [-3,3])
hold on
hline(a0)


```



```

pow_gauss_R = @(mu) 1 - normcdf(norminv(1-a0,0,1) + sqrt(n) .* (mu0 - m
fplot(pow_gauss_R, [-3,3])
hline(a0)
vline(0, '-k')
hold off

```



### Example 4.16

Solution of the system to find the minimal sample size.

```

[51] syms n positive
      solve((n+1)/2 + 1.645 * sqrt(n)/2 - n * 0.6 - 1/2 + 0.842 * sqrt(n*0.24),

ans =
((421*6^(1/2))/250 + 329/40)^2

```

```

[52] eval(ans)

```

```

ans =
152.5210

```

Transform  $H_0 : p_\mu \leq 0.05$  where  $p_\mu = P_\mu(X > 100)$  into  $H'_0 : \mu \leq 85.05$ .


[53]

```
syms mu positive
fzero(@(mu) poissinv(0.95, mu) - 100.5, 50)
```

```
ans =
    85.0571
```

[54]

```
fplot(@(mu) poissinv(0.95, mu), [0, 110])
hline(100)
vline(85.0571)
```




## The t distribution

Figure 4.10. Densities of the t-distributions with 1 (blue), 5 (red), and  $\infty$  degrees of freedom.

[55]

```
fplot(@(x) tpdf(x, 1), [-4, 4]); hold on
fplot(@(x) tpdf(x, 5), [-4, 4]);
fplot(@(x) normpdf(x), [-4, 4]); hold off
```




### Example from Altman

Daily intake in KJ for 11 women.

```
[56] daily_intake = [5260,5470,5640,6180,6390,6515,6805,7515,7515,8230,8770]
```

```
[57] qqplot(daily_intake)
```



Test that  $H_0 : \mu \leq 6000$  against  $H_1 : \mu > 6000$

First with Gauss test assuming  $\sigma = 1200$

```
[58] [H, p, ~, stats] = ztest(daily_intake, 6000, 1200, 'tail', 'right', 'al
```

```
H =
     1
p =
    0.0186
stats =
    2.0829
```

```
[59] (mean(daily_intake) - 6000)/(1200/sqrt(11))
```

```
ans =
    2.0829
```

```
[60] 1-normcdf(2.0829)
```

```
ans =
    0.0186
```

Test significant.

Then use a t-test.

```
[61] [H, p, ~, stats] = ttest(daily_intake, 6000, 'tail', 'right', 'alpha',  
  
H =  
    0  
p =  
    0.0267  
stats =  
    struct with fields:  
  
    tstat: 2.1885  
    df: 10  
    sd: 1.1421e+03
```

```
[62] sqrt(11) * (mean(daily_intake) - 6000)/std(daily_intake)  
  
ans =  
    2.1885
```

```
[63] pval = 1-tcdf(2.1885,10)  
  
pval =  
    0.0267
```

## Example. Sign test

Sign test for the hypothesis  $H_0 : \mu \leq 6000$  where  $\mu$  is now the median.

```
[64] sign(daily_intake - 6000)  
  
ans =  
    -1    -1    -1     1     1     1     1     1     1     1     1
```

```
[65] pval = 1-binocdf(7, 11, 1/2)  
  
pval =  
    0.1133
```

```
[66] [p,H] = signtest(daily_intake, 6000,'tail','right')
```

```
p =  
    0.1133  
H =  
    logical  
    0
```

### Example 4.33 Student t test for paired samples

Example from Dobson. The weights, in kilograms, of twenty men before and after participation in a 'waist loss' program are considered (Egger et al., 1999) We want to know if, on average, they retain a weight loss twelve months after the program.

```
[67] before = [100.8 102.0 105.9 108.0 92.0 116.7 110.2 135.0 123.5 95.0 ...  
105.0 85.0 107.2 80.0 115.1 103.5 82.0 101.5 103.5 93.0]';  
after = [  
97.0 107.5 97.0 108.0 84.0 111.5 102.5 127.5 118.5 94.2 ...  
105.0 82.4 98.2 83.6 115.0 103.0 80.0 101.5 102.6 93.0]';  
XY = [before after]
```

```
XY =  
    100.8000    97.0000  
    102.0000    107.5000  
    105.9000    97.0000  
    108.0000    108.0000  
     92.0000    84.0000  
    116.7000    111.5000  
    110.2000    102.5000  
    135.0000    127.5000  
    123.5000    118.5000  
     95.0000     94.2000  
    105.0000    105.0000  
     85.0000     82.4000  
    107.2000     98.2000  
     80.0000     83.6000  
    115.1000    115.0000  
    103.5000    103.0000  
     82.0000     80.0000  
    101.5000    101.5000  
    103.5000    102.6000  
     93.0000     93.0000
```


```
[68] z = before - after;  
qqplot(z)  
[H, p, ~, stats] = ttest(z, 0, 'tail', 'right')
```

```

H =
    1
p =
    0.0049
stats =
    struct with fields:

    tstat: 2.8734
    df: 19
    sd: 4.1167

```



Assuming normality there is a strong evidence against the null  $H_0 : \Delta \leq 0$  because  $p < 0.01$ . Normality is however questionable.


## Example t-test for independend samples

Example from Maindonald. Consider data from an experiment in which 21 elastic bands were randomly divided into two groups, one of 10 and one of 11. Bands in the first group were immediately tested for the amount that they stretched under a weight of 1.35 kg. The other group were dunked in hot water at 65° C for four minutes, then left at air temperature for ten minutes, and then tested for the amount that they stretched under the same 1.35 kg weight as before.

```
[69] ambient= [254 252 239 240 250 256 267 249 259 269]';
```

```
[mean(ambient) std(ambient)]
[mean(heated) std(heated)]
boxplot([ambient; heated], [ones(10,1); ones(11,1)*2])
```

```
ans =
    253.5000    9.9247
ans =
    244.0909   11.7342
```



Is the difference of means significant (provided that the variances were equal)?

```
[70] [H, p, ~, st] = ttest2(ambient, heated, 'tail', 'both')
```

```
H =
    0
p =
    0.0632
st =
    struct with fields:
        tstat: 1.9730
        df: 19
        sd: 10.9145
```

The difference is not significant (actually there is a borderline evidence against the null).




## Empirical Distribution function

Figure 4.12. The empirical distribution function of a sample of size 25 from the  $N(0, 1)$  distribution and the actual distribution function.

```
[71] clear
      rng(204)
      n = 25;
      x = normrnd(0,1,25,1);
```

```
[72] cdfplot(x)
      hold on
      fplot(@normcdf, [-3,3])
      hold off
```



## Kolmogorov Smirnov test

Data on grades (scale 0-100) on an exam. The a Kolmogorov-Smirnov test is applied for the null  $H_0 : F \sim F_0 = N(75, 10)$ . That is test the null hypothesis that the data comes from a normal distribution with a mean of 75 and a standard deviation of 10. First data are standardized with respect to  $\mu_0 = 75$  and  $\sigma_0 = 10$  and the empirical distribution function is plotted against the standard normal.

```
[3] load examgrades
test1 = grades(:,1);
[mean(test1) std(test1)]
x = (test1-75) ./ 10;
```

```
ans =
    75.0083    8.7202
```

```
[4] [test1 x]
```


```
ans =
    65.0000   -1.0000
    61.0000   -1.4000
    81.0000    0.6000
    88.0000    1.3000
    69.0000   -0.6000
    89.0000    1.4000
    55.0000   -2.0000
    84.0000    0.9000
    86.0000    1.1000
    84.0000    0.9000
    71.0000   -0.4000
    81.0000    0.6000
    84.0000    0.9000
    81.0000    0.6000
    78.0000    0.3000
    67.0000   -0.8000
    96.0000    2.1000
    66.0000   -0.9000
    73.0000   -0.2000
    75.0000         0
    59.0000   -1.6000
    71.0000   -0.4000
    69.0000   -0.6000
    63.0000   -1.2000
    79.0000    0.4000
    76.0000    0.1000
    63.0000   -1.2000
    85.0000    1.0000
    87.0000    1.2000
    88.0000    1.3000
    80.0000    0.5000
    71.0000   -0.4000
    65.0000   -1.0000
    84.0000    0.9000
    71.0000   -0.4000
    75.0000         0
    81.0000    0.6000
```

64.0000	-1.1000
65.0000	-1.0000
84.0000	0.9000
77.0000	0.2000
70.0000	-0.5000
75.0000	0
84.0000	0.9000
75.0000	0
73.0000	-0.2000
92.0000	1.7000
90.0000	1.5000
79.0000	0.4000
80.0000	0.5000
71.0000	-0.4000
73.0000	-0.2000
71.0000	-0.4000
58.0000	-1.7000
79.0000	0.4000
73.0000	-0.2000
64.0000	-1.1000
77.0000	0.2000
82.0000	0.7000
81.0000	0.6000
59.0000	-1.6000
54.0000	-2.1000
82.0000	0.7000
57.0000	-1.8000
79.0000	0.4000
79.0000	0.4000
73.0000	-0.2000
74.0000	-0.1000
82.0000	0.7000
63.0000	-1.2000
64.0000	-1.1000
73.0000	-0.2000
69.0000	-0.6000
87.0000	1.2000
68.0000	-0.7000
81.0000	0.6000
73.0000	-0.2000
83.0000	0.8000
73.0000	-0.2000
80.0000	0.5000
73.0000	-0.2000
73.0000	-0.2000
71.0000	-0.4000
66.0000	-0.9000
78.0000	0.3000
64.0000	-1.1000
74.0000	-0.1000
68.0000	-0.7000

75.0000	0
75.0000	0
80.0000	0.5000
85.0000	1.0000
74.0000	-0.1000
76.0000	0.1000
80.0000	0.5000
77.0000	0.2000
93.0000	1.8000
70.0000	-0.5000
86.0000	1.1000
80.0000	0.5000
81.0000	0.6000
83.0000	0.8000
68.0000	-0.7000
60.0000	-1.5000
85.0000	1.0000
64.0000	-1.1000
74.0000	-0.1000
82.0000	0.7000
81.0000	0.6000
77.0000	0.2000
66.0000	-0.9000
85.0000	1.0000
75.0000	0
81.0000	0.6000
69.0000	-0.6000
60.0000	-1.5000
83.0000	0.8000
72.0000	-0.3000

```
[5] cdfplot(x); hold
     fplot(@(x) normcdf(x), [-3,3], 'r-')
```

Current plot held



```
[6] [h,p,kstat] = kstest(x)
```

```
h =  
    logical  
    0  
p =  
    0.5612  
kstat =  
    0.0707
```

There is no evidence against the null.

## Chi square test

The test statistic is

$$X^2 = \sum_{j=1}^k \frac{(N_j - np_j)^2}{np_j}$$

Same example as before.

Test against standard normal (mean 0, standard deviation 1).

```
[24] [h,p, stats] = chi2gof(x,'cdf',@normcdf)

h =
    0
p =
    0.1153
stats =
    struct with fields:

        chi2stat: 14.2017
           df: 9
        edges: [1x11 double]
           O: [4 5 14 11 29 14 22 14 4 3]
           E: [1x10 double]
```

Test against a general normal with parameters estimated. The degrees of freedom are  $df = k - 1$ . The number of bins is 10 by default.

```
[34] [h,p, stats] = chi2gof(test1)

h =
    0
p =
    0.0626
stats =
    struct with fields:

        chi2stat: 10.4832
           df: 5
        edges: [1x9 double]
           O: [9 14 11 29 14 22 14 7]
           E: [8.8927 11.2029 17.6673 22.1913 22.2013 17.6913 11.2283
8.9249]
```

Degrees of freedom are  $df = k - 1 - 2$  where 2 is the number of estimated parameters. Here  $k = 8$ .

More information with `help chi2gof`

## Likelihood ratio test

A simple example: assume that  $X \sim \text{Bin}(100, p)$  and that we want to test  $H_0 : p = 1/2$  against  $H_1 : p \neq 1/2$ . The likelihood ratio statistic is

$$\lambda(X) = \frac{L(\hat{p}, X)}{L(p_0, X)}.$$


We reject for large values of the statistic. Below you see the value of the statistic for  $x = 60$ .

```
[77] n = 100; x = 60; phat = x/n; p0 = 1/2;  
lambda = binopdf(x, n, phat) / binopdf(x, n, p0)
```

```
lambda =  
    7.4899
```

```
[78] fplot(@(p) binopdf(x, n, p), [0,1])  
vline(0.5, '-')
```

```
hline(binopdf(x, n, p0))
```




## The logarithm of the likelihood ratio test

Consider twice the log of the likelihood ratio statistic

$$2 \log L(\hat{p})/L(p_0) = 2\{\ell(\hat{p}) - \ell(p_0)\}$$

We plot it as a function of  $p_0$  for a given  $x = 60$ .

```
[79] fplot(@(p0) 2 * log( binopdf(60, 100, 0.6) ./ binopdf(60, 100, p0)), [0
vline(0.5, '-')
```



The function shows the value of the statistic that increases as we move away from the MLE  $\hat{p} = 0.6$ .

## A simulation for the sampling distribution of the likelihood ratio

We simulate many samples of size  $n = 500$  under the hypothesis  $p = 1/2$ , computing the statistic  $2 \log \lambda$ . We verify that the empirical distribution is approximately  $\chi_1^2$ .

```
[80] B = 1000;
p0 = 1/2;
n = 500;
x = zeros(1,B);
LR = zeros(1,B);
for b = 1:B
```



```

    ph = x(b)/n;
    LR(b) = 2 * log( binopdf(x(b), n, ph) ./ binopdf(x(b), n, p0));
end

histogram(LR, 40, 'Normalization', 'pdf')
hold on
fplot(@(x) chi2pdf(x,1), [0,10])

```

